

Technical Appendix

Details of Meta-sketch Operations

Algorithm 1 describes details about operations of the meta-sketch, including broadcast, dimensions conversion processes, and three forms of \ominus . It is important to emphasize that all the operations are performed on one stream item, so that a parallelized version can easily be implemented by adding an additional dimension.

Algorithm 1: Details of Meta-sketch Operations

```

1 Operation Store ( $e_i, M$ ):
2    $z_i, r_i \leftarrow \mathcal{F}_E(e_i); a_i \leftarrow \mathcal{F}_{Sa}(r_i);$ 
3    $a_i \leftarrow \text{changeShape}(a_i, \mathbb{R}^{d_1 \times d_2}, \mathbb{R}^{d_1 \times 1 \times d_2})$ 
4    $z_i \leftarrow \text{changeShape}(z_i, \mathbb{R}^{l_z}, \mathbb{R}^{d_1 \times l_z \times 1})$ 
5    $M \leftarrow M + z_i a_i;$ 
6 Operation Delete ( $e_i, M$ ):
7    $z_i, r_i \leftarrow \mathcal{F}_E(e_i); a_i \leftarrow \mathcal{F}_{Sa}(r_i);$ 
8    $a_i \leftarrow \text{changeShape}(a_i, \mathbb{R}^{d_1 \times d_2}, \mathbb{R}^{d_1 \times 1 \times d_2})$ 
9    $z_i \leftarrow \text{changeShape}(z_i, \mathbb{R}^{l_z}, \mathbb{R}^{d_1 \times l_z \times 1})$ 
10   $M \leftarrow M - z_i a_i;$ 
11 Operation Query ( $x_i, M, N$ ):
12   $z_i, r_i \leftarrow \mathcal{F}_E(x_i); a_i \leftarrow \mathcal{F}_{Sa}(r_i);$ 
13   $\hat{f}_i \leftarrow \mathcal{F}_{dec}(\{M \ominus a_i\}, z_i, N);$ 
14  return  $\hat{f}_i;$ 
15 Module Embedding ( $x_i$ ):
16   $z_i \leftarrow g_{emb}(e_i); r_i \leftarrow g_{add}(z_i);$ 
17  return  $z_i, r_i$ 
18 Module SparseAddress ( $r_i$ ):
19   $r_i \leftarrow \text{changeShape}(r_i, \mathbb{R}^{l_r}, \mathbb{R}^{d_1 \times 1 \times l_r});$ 
20   $\hat{a}_i \leftarrow r_i A;$ 
21   $\hat{a}_i \leftarrow \text{changeShape}(\hat{a}_i, \mathbb{R}^{d_1 \times 1 \times d_2}, \mathbb{R}^{d_1 \times d_2});$ 
22   $a_i \leftarrow \text{SparseMax}(\hat{a}_i, \text{dim} = -1)$ 
23  return  $a_i$ 
24 Module Decoding ( $\{M \ominus a_i\}, z_i, N$ ):
25   $m_i \leftarrow \text{basicRead}(M, a_i)$ 
26   $i_1, i_2 \leftarrow \text{advancedRead}(m_i, z_i)$ 
27   $\text{info} \leftarrow \text{concatenate}(m_i.\text{flatten}(), i_1, i_2, N)$ 
28   $\hat{f} \leftarrow g_{dec}(\text{info})$ 
29  return  $\hat{f}$ 
30 Function changeShape ( $\text{vector}, \mathbb{R}^n, \mathbb{R}^m$ ):
31  change vector's shape from  $\mathbb{R}^n$  to  $\mathbb{R}^m$ 
32  return vector
33 Function basicRead ( $M, a_i$ ):
34   $a_i \leftarrow \text{changeShape}(a_i, \mathbb{R}^{d_1 \times d_2}, \mathbb{R}^{d_1 \times d_2 \times 1})$ 
35   $m_i \leftarrow M a_i$ 
36   $m_i \leftarrow \text{changeShape}(m_i, \mathbb{R}^{d_1 \times l_z \times 1}, \mathbb{R}^{d_1 \times l_z})$ 
37  return  $m_i$ 
38 Function advancedRead ( $m_i, z_i$ ):
39   $z_i \leftarrow \text{changeShape}(z_i, \mathbb{R}^{l_z}, \mathbb{R}^{d_1 \times l_z})$ 
40   $i_1 \leftarrow m_i.\text{min}(\text{dim} = -1)$ 
41   $z_i^1 \leftarrow \text{where}(z_i > \epsilon, z_i, \epsilon)$ 
42   $z_i^2 \leftarrow \text{where}(z_i < \epsilon, \text{MAX}, 0)$ 
43   $i_2 = [(m_i + z_i^2)/z_i^1].\text{min}(\text{dim} = -1)$ 
44  return  $i_1, i_2$ 

```

Details of meta-task generation

The detailed algorithms for generating basic/adaptive meta-tasks are shown in Algorithm 2 and Algorithm 3, respectively.

Algorithm 2: Generating a Basic Meta-task

```

Data: Item pool  $I$ ; Distribution pool  $P$ ; Frequency mean range  $L$ ;
Result: a meta-task  $t_i$ ;
1 Sample an item size  $n_i$  from  $[1, |I|]$ ;
2 Sample a frequency mean  $\bar{f}$  from  $L$ ;
3 Sample an subset  $\{x_1^{(i)}, \dots, x_{n_i}^{(i)}\}$  of  $I$  with size  $n_i$ ;
4 Sample a instance  $p^{(i)} \sim P$ ;
5 for  $x_j^{(i)} \in \{x_1^{(i)}, \dots, x_{n_i}^{(i)}\}$  do
6   Sample  $p_j^{(i)} \sim p^{(i)}$  and  $f_j^{(i)} \leftarrow \lceil n_i \times \bar{f} \times p_j^{(i)} \rceil$ ;
7   add  $x_j^{(i)}$  to the  $t_i$ 's store set ( $s_i$ ) with  $f_j^{(i)}$  times;
8   add  $(x_j^{(i)}, f_j^{(i)})$  to  $t_i$ 's query set ( $q_i$ );
9 end

```

Algorithm 3: Generating an Adaptive Meta-task

```

Data: Item pool  $I$ ; Real frequency distribution  $p$ ; Frequency mean range  $L$ ;
Result: a meta-task  $t_i$ ;
1 Sample an item size  $n_i$  from  $[1, |I|]$ ;
2 Sample a frequency mean  $\bar{f}$  from  $L$ ;
3 Sample an subset  $\{x_1^{(i)}, \dots, x_{n_i}^{(i)}\}$  of  $I$  with size  $n_i$ ;
4 for  $x_j^{(i)} \in \{x_1^{(i)}, \dots, x_{n_i}^{(i)}\}$  do
5   Sample  $p_j \sim p$  and  $f_j^{(i)} \leftarrow \lceil n_i \times \bar{f} \times p_j \rceil$ ; // The
   correspondence between items and
   frequencies is changed.
6   add  $x_j^{(i)}$  to the  $t_i$ 's store set ( $s_i$ ) with  $f_j^{(i)}$  times;
7   add  $(x_j^{(i)}, f_j^{(i)})$  to  $t_i$ 's query set ( $q_i$ );
8 end

```

Hyper-Parameters

We did not deliberately tune the parameters of the meta-sketch. We just followed the setting about conventional NN to choose parameters by balancing the sketching ability and training efficiency. Table 1 shows all hyper-parameters that are considered (best parameters are bolded).

Learning rate of MS	$\{1e-3, 5e-4, \mathbf{1e-4}, 5e-5\}$
Hidden Size of g_{emb}	$\{64, \mathbf{128}, 256\}$
Hidden Size of g_{add}	$\{24, \mathbf{48}, 64, 64\}$
Hidden Size of g_{dec}	$\{128, \mathbf{256}, 512\}$
$d_2 : l_r$	$\{5 : 1, \mathbf{5 : 2}, 5 : 4, 5 : 6\}$
d_1	$\{1, \mathbf{2}, 3\}$

Table 1: Hyper-parameters Considered

Ablation Study

As shown in Figure 1, we conduct ablation studies to evaluate some key techniques of the meta-sketch. In all comparisons,

the settings follow experiment Section ($n = 5K, B = 9KB$, Word-query), as shown in Table 2. The comparison between Base and Abl 1 shows the effectiveness of the optimizations on operation \ominus . The comparison between Base and Abl 2 shows improvement with the address network, especially for the later stages of meta-sketch training. It should be emphasized that embedding vector will pass a *Relu* activation before the output of the g_{emb} , which allows the model to control the sparsity of embedding vectors easily. In the comparison between Base and Abl 3, we can see the effectiveness of the *Relu*.

	Base	Abl1	Abl2	Abl3
\ominus	yes	no	yes	yes
g_{add}	yes	yes	no	yes
<i>Relu</i>	yes	yes	yes	no

Table 2: Settings for the Ablation Study




Figure 1: Ablation Study

The Default Settings and Discussion for $M(A)$

The default parameters of $M(A)$ under different budgets are shown in Table 3. Note that we first set the size of M , and then obtain the size of the compressed A according to the ratio of $l_z : l_r \approx 5 : 1$.

We further discuss the effect of the setting for M . Figure 2 shows the effect of different settings in Table 4 on the training of the meta-sketch under a fixed 9KB budget. All competitors follow the same training setting ($n = 5K, B = 9KB$, Word-query). For d_1 , we set it in the range of 1 to 2, similar to the setting of the number of hash functions in traditional




Figure 2: AAE w.r.t. Different Settings of M

sketches. Figure 2 shows that when $d_1 = 2$, the model yields a better result. For the settings of d_2 and l_z , it shows that the model yields a better result when the ratio of d_2/l_z is around 2. Thus, we set the default parameters for our experiments under the premise of $d_1 = 2$, and $d_2/l_z \approx 2$.

In addition, we can see an inappropriate setting may harm the stability in the early training phase, leading to non-convergence. For example, we can observe that with a additional dimension M corresponds to the better training stability, in the comparison between settings 1 and 4. The comparison of setting 1, setting 2, and setting 3 also shows a reasonable large l_z is beneficial to the stability of the meta-sketch.

B	5KB	7KB	9KB	11KB	13KB	15KB	17KB
d_1	2	2	2	2	2	2	2
d_2	40	45	50	61	61	64	70
l_z	16	20	23	23	27	30	31
l_r	4	4	5	5	5	6	6

Table 3: Default Size of $M(A)$ for Different Space Budgets

Setting	1	2	3	4
d_1	2	2	2	1
l_z	23	12	6	34
d_2	50	100	200	68

Table 4: Settings of M

Supplementary Experiments

Figure 3 and 4 show the AAEs of different competitors in the experiments of basic meta-sketch Section under different space budgets (B) and different item sizes (n), respectively. Figure 5 compares the ARE of advanced MS and LS under dynamic streaming scenarios in advanced meta-sketch Section.




Figure 3: AAE w.r.t. B

The Parameters of Skewed Distributions

Table 5 shows the parameter settings of the three distributions in analysis Section with different skewness levels. Here, the level of skewness is a relative concept under each type of distribution. We can convert all distributions to a zipf form, i.e. sorting n items on a descending order of $\frac{f}{N}$. Afterwards,

	Level1	Level2	Level3	Level4
Zipf	$\alpha = 1.0$	$\alpha = 0.8$	$\alpha = 0.6$	$\alpha = 0.4$
Triangular	$k = -1/128$	$k = -1/64$	$k = -1/32$	$k = -1/16$
Uniform	$a=0, b=10000$	$a=1250, b=8750$	$a=2500, b=7500$	$a=3750, b=6250$

Table 5: The Parameters of Skewed Distributions

Device	Write Latency		Query Latency		Write Throughput(QPS)		Query ThroughPut(QPS)	
	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU
Meta-sketch	0.80ms	1.32ms	1.57ms	2.25ms	166.69k	7142.86k	135.14k	4063.39k
CM-sketch	0.27ms	-	0.25ms	-	4.80k	-	4.86k	-

Table 6: Latency and Throughput of Meta-sketch




Figure 4: AAE w.r.t. n




Figure 5: LS vs. MS on ARE

the level of skewness can be measured by the slope from the first to last positions of the ordering.

Latency and Throughput of Meta-sketch

We evaluate the write/query latency and throughput of the meta-sketch and the CM-sketch¹ under the same setting. We use a single write/query operation for the testing of the latency and a batch of 10K write/query operations for the testing of the throughput. As shown in Table 6, the latency of write/query operations of the meta-sketch is slightly higher than that of the CM-sketch. But with the parallel algebraic operations of NNs, meta-sketch can have a significantly higher throughput, e.g., in GPU environment. For example, the writing throughput of meta-sketch is around 30/1400

¹The implementation of the CM-sketch is from package pyprobables, which is a definitive python library for probabilistic data structures (<https://pyprobables.readthedocs.io/en/latest/code.html>).

times higher than that of the CM-sketch when deployed on CPU/GPU. Similar observations are drawn on query operations.